

A Data Revolution in the Cognitive Sciences

John V. McDonnell (*moderator*) and Todd M. Gureckis

New York University Department of Psychology

Johan Bolen

Indiana University School of Informatics and Computing

Georg Langs

Medical University of Vienna

Winter Mason, Meeyoung Cha, Krishna Gummadi, Farshad Kooti, and Haeryun Yang

Howe School of Technology Management, Stevens Institute

John Myles White

Princeton University Department of Psychology

Keywords: Machine Learning; Cognitive Neuroscience; Online Experiments; Social Network Analysis; Text Analysis; Methods

In the 1980s, the broad availability of PCs fueled a renaissance in cognitive science, lowering computational barriers to the development of sophisticated computational models of cognition and changing the way we run experiments as well as the tools we use for analysis. As technology continues to advance, the advent of big data is now leading to methodological advances on various fronts, from online experimentation to the mining of large datasets. These new opportunities also come with new challenges, as scientists must overcome the challenges these methods bring. The goal of this symposium is to highlight how researchers are making use of these new tools and the techniques they have used to solve the challenges they represent.

A major theme of the symposium will be the way in which these new techniques represent more than just a scaling up of previous efforts: Researchers are now able to ask new questions in different ways. For example, collective behavior has long been of interest to cognitive scientists but was difficult to study directly. Now that scientists have access to anonymized data from massive social networks, these interactions can be studied in the environment. Johan Bolen and Winter Mason will be discussing work they have done using this newly available data source.

Another novel source of data are online labor marketplaces such as Amazon's Mechanical Turk (AMT) service. These allow researchers to perform controlled studies involving large numbers of participants in short periods of time. This allows researchers to run experiments requiring far more data than in the past, such as brief one-shot studies with thousands of participants or learning task optimization. John Myles White and John McDonnell will talk about their work running studies online with AMT.

Finally, cognitive neuroscience has also been a major beneficiary of the data revolution. The field has amassed im-

mense quantities of scanning data from fMRI and other imaging studies. Analysis of these data poses challenges in terms of statistical rigor (given the problem of multiple comparisons on a massive scale) and managing immense quantities of data. Georg Langs will speak about his work using cutting edge machine learning techniques to derive new insights from imaging data.

Big data for computational social science: Social networks and sentiment

Johan Bollen Twitter and Facebook now have more than one billion users combined. These social media users produce hundreds of millions of messages each day that provide a unique window into our collective thoughts, feelings, experiences, and observations. Social media data is uniquely valuable, not merely because of its scale, coverage, and ability to gauge ephemeral personal conditions, but because it provides a record of long-term social ties and demographic features. The resulting data enables the study of a variety of socio-economic phenomena at the intersection of psychology, sociology, linguistics, biology, and computer science.

Our research is particularly focused on the role that human mood and sentiment play in collective intelligence and decision-making. At the individual level our decision-making is strongly influenced by emotions. This may be even more the case at a societal level where social relations and mass communication amplify the production and propagation of emotive information and misinformation.

In past work we have performed large-scale analysis of collective mood states using computationally extended versions of existing psychometric models. The results provide a quantitative assessment of the fluctuations of societal well-being and mood over time that can be leveraged to study a variety of sociological phenomena such as mood homophily, mood contagion, meme propagation, and the effect of public mood states on economic and financial indicators. More recent work has focused on the development of computational models of the relations between social mood and socio-economic

indicators, as well as the analysis of collective knowledge representations and their role in the propagation of misinformation and political bifurcation processes.

Machine Learning Approaches in Neuroimaging and Medical Image Analysis.

Georg Langs The computational analysis of biomedical imaging data has become central in medical research and clinical application. The talk will highlight two recent developments that might shape research in this area in the near future. Both are connected to the impact of the emerging ability to process large data, and the approaches that become relevant. The advent of the capability to process large data—hundreds of terabytes—has changed paradigms in how we approach classical pattern recognition problems in this domain. Medical research fields that rely on exploratory approaches together with large and complex data to study and understand physiological processes are increasingly shaped by a tight interaction among researchers in machine learning and medical sciences. In the talk I will discuss these directions, highlight open issues, and illustrate them with examples from medical image retrieval, and neuroimaging research and clinical application.

The emergence of social conventions in social networks

Winter Mason, Meeyoung Cha, Krishna Gummadi, Farshad Kooti, and Haeryun Yang Social conventions and norms are a powerful and ubiquitous guide for behavior. However, the way in which conventions emerge in communities is not very well understood, largely due to a paucity of available data. In this study, we leverage the widespread use of the micro-blogging platform Twitter and focus on competing conventions for attributing reposts to the original source. We analyze over 1.7 billion “tweets” from 54 million users, from the very first tweet in 2006 up to September 2009, and observe how the conventions emerged and spread through the network of Twitter users. We observe that initially the most successful conventions were borrowed from natural language (“via” and “retweeting”), but eventually a community-specific convention came to dominate (“RT”). We see that some conventions are used by divergent groups of people, while others are abandoned in favor of more efficient expressions. We also observe the failure of some conventions despite higher efficiency (i.e., fewer characters) and explicit endorsement of their adoption. Our results suggest that there are some features that encourage the adoption of one convention over another, but that there is still significant inherent unpredictability in what convention will come to dominate.

Context and decision making in a massive online experiment

John Myles White Online labor markets, like Amazon’s Mechanical Turk (AMT), offer psychologists many opportunities. The convenience of the virtual lab provided by

AMT has already won over many psychologists, but transitioning research from the lab to the web browser offers other benefits. AMT allows psychologists to deploy experiments that are fully automated: The recruitment, instruction, core experiment and debriefing periods can be identical for all subjects, which largely removes the possibility that undocumented components of an experiment might exert substantive influences on the final results (Doyen et al., 2012). Moreover, the ease of recruiting large numbers of subjects for tasks provides important increases in statistical power, which can allay concerns about false positive psychology (Simmons et al., 2012). But I will argue that the primary value of AMT is not purely methodological: the strength of the virtual lab is that it allows psychologists to pursue large-scale between-subjects designs in which a large number of subjects perform a very small number of trials. This type of research is often eschewed because of the prohibitive cost of recruiting hundreds or thousands of subjects in the lab, but previous research (e.g. Gneezy et al. 2006) suggests that studies of decision-making can be powerfully influenced by contextual cues. Our recent work has found evidence that the effects of context on decisions may be even more problematic than previously believed: in our studies of decision-making, we find that the effects of ostensibly innocuous local context can shift the basic qualitative results of experiments.

Can online data be trusted? Learning tasks on Amazon’s Mechanical Turk.

John V. McDonnell, Todd M. Gureckis Amazon Mechanical Turk (AMT) has attracted attention from experimental psychologists interested in gathering human subject data more efficiently. However, relative to traditional laboratory studies, many aspects of the testing environment are not under the experimenter’s control. We have empirically evaluated the fidelity of AMT for use in cognitive behavioral experiments. These types of experiment differ from simple surveys in that they require multiple trials with sustained attention from participants. Our initial attempts to replicate the classic Shepard, Hovland, and Jenkins (1961) task were only partially successful. However, after systematically studying the effects of compensation and validation we were able to more closely match previous in-lab findings. Specifically, we found that compensation altered the rate of sign-ups, but not the quality of the data. Conversely, we found that using a strict measure to ensure that participants had understood instructions resulted in a considerable improvement in performance and convergence with in-lab findings.

References

- Doyen, S., Klein, O., Pichon, C. L., & Cleermans, A. (2012). Behavioral priming: It’s all in the mind, but whose mind? *PloS one*, 7, e29081.
- Gneezy, U., List, J. A., & Wu, G. (2006). The uncertainty effect: When a risky prospect is valued less than its worst

possible outcome. *The Quarterly Journal of Economics*, 121, 1283–1309.

Shepard, R., Hovland, C., & Jenkins, H. (1961). Learning and memorization of classifications. *Psychological Monographs*.

Simmons, J. P., Nelson, L. D., & Simonsohn, U. (2011). False-positive psychology undisclosed flexibility in data collection and analysis allows presenting anything as significant. *Psychological Science*, 22, 1359–1366.